

What the current run established—and what it did not

The model generated every response and action request naturally; no successful compromise was scripted into the data.

8

clean trajectories

8

injected trajectories resisted

0

unsafe actions executed

Supported by the completed run

- ✓ The live 2-hop and 3-hop pipeline completed successfully.
- ✓ Exact model records and all-layer activations were preserved.
- ✓ Group-held-out construction-signal scoring was completed.



Evidence required for the next claim

- Expand controlled domains and injection wordings.
- Observe executed, propagated, blocked, and resisted outcomes.
- Then estimate behavioral AUROC and delegation-depth effects.

Current conclusion: the end-to-end measurement loop works, but behavioral-compromise detection remains an open empirical question.